

COMP 3220
Homework 6
Due: Monday, November 3rd, 2014, 11:59PM
(Please submit using Canvas)

This assignment is worth 100 points.

In this assignment, you are required to use Java competition and cooperation synchronization features that facilitate programming concurrent tasks (see section 13.7 in your textbook). In class, we used the *Producer-Consumer* problem as an example to demonstrate Java's concurrency control features. In this assignment, you will develop a solution for a variant of the problem.

Problem Definition: A shared memory location (e.g., a buffer) can be concurrently read by any number of tasks, but when a task must write to the shared memory location, it must have exclusive access. Write a Java program for coordinating multiple concurrent tasks (Java *Threads*), each one representing either a reader or writer. **You need to use Java threads and synchronized methods to implement a monitor to coordinate the readers and writers.** In the *main()* function of your program, please make sure to create multiple readers and writers and then allow them to access the buffer concurrently.

Please submit a ZIP (compressed) package, which includes (1) the source code that follows the style guidelines, (2) a text document that provides a trace of a single run showing (using print statements) instantiation of multiple readers and writers and demonstrating that writers and readers are accessing the buffer in accordance with the synchronization constraints (e.g., many readers can read simultaneously, but a writer must have exclusive access) and (3) instructions to compile, execute, and test the program.

PROGRAMMING STYLE GUIDELINES

These notes provide guidelines for internal program documentation and style. Although they are intended for student programmers, style skills will carry over into professional life after school. In fact, most professional programmers would consider these standards to be less than minimum requirements for commercial-quality software documentation. While geared mainly for a block structured language such as Pascal or C, most of the points discussed are applicable to any programming language. We will not address the ``how" of program specification, design, or testing; these topics are the domain of the Computer Science curriculum.

The essence of good programming style is communication. Good style in programming is roughly as difficult to learn as good style in English. In both cases, the document has no value if it does not convey its meaning to the reader. Any program that will be used must be maintained by somebody - and that somebody must be able to understand the code by reading it. Any program that needs debugging will be easier to debug if the creator carefully explains what's going on.

Within the program text, programmers have three primary tools for communicating their intentions: comments (explanation for the program); clear variable names, constants, expressions and subroutines (the words of the program itself); and white space and alignment (the organization of the words in the program). Each of these aspects aid communication between the program writer and the program reader (which includes the program writer at debug time - so you as a program writer have a stake in good style too!).

General Guidelines

The most useful things to know about program documentation are ``what" and ``when". In general, you should include comments explaining what every subroutine does, what every variable does, and an explanation for every tricky expression or section of code. But, there is much to good style beyond comments, as we shall see. ``When" is easy - comment it before you write it. Whenever you declare a variable, include a comment. Whenever you start to write a subroutine, first write a comment explaining it. That will help to make clear in your mind what you are about to do. If you find the explanation for a subroutine difficult to write, its a sure sign that you have done a poor job of structuring the program. Avoid obscure programming language constructs and reliance on language-specific precedence rules. It is often better to force precedence by use of parentheses since this leaves no doubt as to meaning. In general, if you had to look up some rule or definition, your readers most likely will too. Whenever you do write a difficult expression, or some other tricky business, if you found it difficult to write you should expect that it will be difficult to understand. So, add a comment.

The most important rule of style is: be consistent! If you adopt some method for variable naming or indenting, stick to it throughout your program.

The Program Header

All programs should include, at or near the beginning of the program, a comment block with at least the following information: the programmer's name and appropriate identification, the date of writing (and possibly dates and purpose of major modifications), the name of the program, the computer and operating system for which it was written. Include information about compiling the program (i.e., the compile line commands). Also include a ``usage" line showing how the program is called and the various command line arguments. The bulk of the program header will often be a description of the program. This serves as a roadmap that others can follow in order to understand the program. At a minimum, include an explanation of how the program works, and the major data structures and algorithms used. If you have more than one module, they should each contain a program header block describing the workings of that module. If you know about any bugs in your program, it is a good idea to include a comment in the program header explaining what they are. You might also want to keep a running ``TODO" list in the program header comment. If you keep a header block ``template" available, you can easily copy this template into your program when you begin. An example program header block in Pascal might look as follows:

```
(*****  
PROGRAM NAME - threshold  
PROGRAMMER - A. Hacker, 001-00-2001  
USAGE - threshold <image-file> <num-rows> <num-cols> <thresh-value>  
        where <image-file> contains an image array of size  
        <num-rows> X <num-cols>, and <thresh-value> is the desired  
        pixel threshold value.  
COMPILE - vp -o threshold threshold.p utils.p  
SYSTEM - VP PASCAL V1.05a on a MAC II running A/UX V2.0  
DATE - Started 11/27/1990  
        Phase I completed 11/30/1990  
BUGS - Program crashes if the number of bytes in <image-file>  
        is less than <num-rows> X <num-cols>.  
DESCRIPTION - This program counts the number of pixels in an  
        image (represented by an array) whose value is  
        greater than a specified threshold.  
*****)
```

Subroutine Documentation

The subroutines of a program (or module) define the structure of that program. Good program structure can be viewed as a business hierarchy: the president (main routine) delegates tasks to the next level of workers, who in turn delegate subtasks to their underlings. Each subroutine should have one clearly defined task. Writing the comment describing that task before writing the subroutine can help to keep you straight here. Subroutines should be short enough so that their full meaning can be grasped as a unit by

the reader. Normally, one page of code represents the limit of intellectual complexity that a reader can grasp.

Every subroutine should be preceded by a comment describing the purpose and use of that subroutine. Similar to the program header, each subroutine should have a comment block with standard sections such as subroutine name, parameters, description, and calling routines. Some programmers like to surround the entire comment block by asterisks to set it off. I don't recommend this practice, nor putting comment braces at the beginning and end of each line in the comment block. While doing so will make the program more attractive, this practice also makes it difficult to modify the comment block - and your documentation should be easy to maintain! Speaking of maintenance, if you decide to modify a subroutine during program development, be sure to modify the comments as well to keep them accurate.

Below is an example of a subroutine comment header block. Note that this example indicates in the header the purpose of each parameter to the routine, and what is returned.

```
(*****  
Function CountBigPix  
PURPOSE: counts and returns the number of pixels  
         in image (of size rows X cols) with value greater than thresh.  
PARAMETERS:  
    image: pixel array containing the image data  
    rows, cols: the number of rows and columns for image  
    thresh: the pixel threshold value -  
            count pixels greater than this  
CALLED BY: main  
SIDE EFFECTS: This function changes no global variables or  
              parameters  
*****)  
function CountBigPix(image : PixArray; rows, cols : integer;  
                    thresh : PixVal) : integer
```

One of the greatest sources of bugs is in the interface between subroutines. It is a common mistake to have type mismatches between the formal parameters and the actual parameters. It is also common to mis-order parameters, or to leave some out. You should always be scrupulous in checking for such problems when writing your program, and you should always check these early on when debugging. One way to minimize the problem is to keep the number of parameters to any one subroutine at a minimum. Seven or more parameters for a subroutine is probably too many - having this many parameters is a good indication that you have structured your program in the wrong way, or that you need to reconsider the global flow of information within your program. One way to reduce the number of parameters is to group logically connected data into a single parameter through the use of records and structures. For example, if you have a program that uses coordinate points, it probably makes more sense to group the coordinate's x and y values into a single record, instead of passing two separate pieces of information to a subroutine. It is also a good idea to adopt a standard ordering, such as all parameters supplying

information to the subroutine appearing before parameters returning information from the subroutine.

Limit the number of parameters passed by reference (var parameters in Pascal, pointers passed in C). There are only two ways that a subroutine can screw up a variable in another subroutine: global variables and variables passed such that they can be changed. The problem in both cases is that you increase the chance that something will be changed in a way that you didn't expect. By definition (since you didn't expect it), this is hard to debug. If your subroutine does change a global variable or formal parameter, you should document this in the subroutine header comment.

There are two classes of subroutines: functions and procedures. Functions should affect precisely one value - the function's return value. This means that functions should change no global variables, nor should they have parameters passed by reference (sometimes you use call-by-reference to avoid duplicating space, but if you do this in a function, be careful not to change the parameter's value). If more than one value is to be returned (or a global variable is changed), use a procedure. Following this policy guarantees that subroutines declared as functions are relatively safe - they can return the wrong value, but they can't cause any other damage. In C all subroutines are technically functions. Use the void type to indicate "procedure-like" (i.e., dangerous) subroutines. You may have noticed that many C standard functions such as read and write are really procedures that appear to violate this rule. It is customary in C for procedures to return a value that has a special meaning: it indicates the status of the procedure's operation - i.e., whether the operation was successful or not.

Indentation and White Space

You should notice that in the code examples I take care to skip spaces or line things up in a special way. Adding "white space" in this way makes your program much more readable. At least two blank lines should separate your subroutines. It is good to put additional blank lines at appropriate places within subroutines to separate major sections, such as to separate declarations from the beginning of the statements. You can also use blank lines to separate logical sections of the subroutine code.

A new level of indentation should be used at every level of statement nesting in your program. The amount of indentation is up to the programmer - there are many different styles. Commonly, programmers will indent a fixed number of spaces at each level. More spaces make the indentation clearer, but they also limit the amount of information on a line after several levels of indentation. The minimum number of spaces at each indentation should be 3, and many programmers use a tab mark (typically 8 spaces). The other common style of indentation is based on the keyword. For example, a for loop would indent one more character than an if statement.

Where should the begin go? Some programmers put it at the end of the control statement, others put it on the next line. Here are some examples.

```

if x < 0 then
  x := 0 (* 3 spaces fixed indentation style *)
else begin
  while x > 0 do begin (* "begin" on same line as "while" *)
    P(x);
    x := x - 1
  end (* While *)
end; (* else *)

```

```

if x < 0 then
  x := 0 (* indentation matches keyword length style *)
else
  begin (* begin on a new line *)
    while x > 0 do (* Indent some more after begin *)
      begin
        P(x);
        x := x - 1
      end; (* while x > 0 *)
      y := y + 1
    end; (* else *)
  end

```

Nested if-then-else statements are often treated differently, to avoid too much indentation.

```

if <expression1> then begin
  <statements>
end
else if <expression2> then begin
  <statements>
end
else begin
  <statements>
end;

```

Often there are ``pretty-printing" utilities available to programmers that automatically indent programs. For C programmers under UNIX, both the utility `cb` and the EMACS text editor provide automatic indenting.

You should not put more than one statement on the same line. Also, it is good practice to keep lines to less than 80 characters. Lines longer than this should be broken into multiple lines, and indented so as to line things up in a logical way, such as:

```

average := (image[x-1][y-1] + image[x][y-1] + image[x+1][y-1] +
            image[x-1][y] + image[x][y] + image[x+1][y] +
            image[x-1][y+1] + image[x][y+1] + image[x+1][y+1]) / 9;

```

Identifiers

Pick identifiers (i.e., names for variables, functions, records, etc.) that communicate the named object's purpose to the reader. You also want to pick identifiers in such a way that a mistake when typing will not likely result in another identifier used within your program. In this way, a typo will be caught as an error by the compiler (undeclared variable) instead of resulting in a difficult to find "logic" error.

In general, one character identifiers are bad since they violate both of these rules. The variable names `o` and `p` are not usually very informative, and a mistyped `o` can easily become a `p`. However, in two cases it is traditional to use one letter variable names. In certain mathematical expressions it is appropriate to use traditional variable names such as `x` or `y`. It is also appropriate to use `i`, `j` and `k` for loop variables and array indices - but for no other purpose!

Many programming languages (such as Pascal and C) allow for arbitrarily long identifiers. Be careful here, since some compilers will only use the first six or eight characters to distinguish between names.

There are a number of rules for good use of variables. An easy one to follow is: never change the value of a loop variable within that loop. In other words, don't do the following:

```
for i = 1 to 10 do
  i := i + 1;
```

Pascal compilers are supposed to flag such abuse as an error, but not all do.

This is an instance of a general principle for variable use: don't overload the logical meaning of a variable. By using `i` as the loop variable, the programmer is signaling to the reader something about `i`. By assigning to the variable, a second meaning is introduced. The program is easier to understand if a second variable is assigned to and used.

Use variables at the smallest possible level of scope. One implication of this rule is that you should minimize the use of global variables. There is a place in good programming style for global variables, and that is for a body of knowledge that will be acted on by many sections of the program, and which is in some sense the major essence of that program. Don't use global variables as a convenient means to communicate between two subroutines.

Be consistent about the use of local variables with the same name throughout the program. If `i` is used as a loop variable in one subroutine, don't use it as something else in another subroutine.

Depending on the programming language, the compiler may or may not distinguish upper and lower case letters. However, a reader certainly can distinguish the difference, and mixing cases correctly can give the reader additional clues about a name. You should adopt a standard usage for case. One suggestion is to have all constants entirely in upper case, all created types (typedefs, records, structures) with only the first letter upper case,

and all variables entirely in lower case. It is also informative if you use common suffixes and prefixes for related variables, such as `inFile`, `outFile`, `errorFile`. Beware that 0 and 1 look a lot like O and I.

Every variable in your program should have a comment describing its use. Typically, this comment appears at the end of the line containing the variable declaration. You should avoid declaring more than one variable on the same line unless all variables on that line have very similar use. Here is an example of variable declarations in Pascal.

```
var
  image : PixArray;    (* The image buffer *)
  rows, cols : integer; (* Number of rows and columns *)
  i, j : integer;      (* Current row and column *)
  thresh : integer;    (* Threshold: count bigger pixels *)
  BigPixCount : integer; (* Total number of pixels in image
                        array greater than thresh *)
```

Since Pascal and ANSI C subroutine definitions do not have a section for formal parameter definitions other than the function header definition line, you need to find another place to document the purpose of the formal parameters. A good idea is to include this in the subroutine header comment, as in the example above.

Constants and Enumerated Types

When you type a mathematical expression such as `val = 7*x + 2`, the numbers 7 and 2 are truly numbers. However, most programmers only rarely use true numbers. More likely, they are dealing with numbers that are logical quantities, such as the size of a disk block (1024) or Monday (day 1 in the week). Of great importance (but often neglected by weaker programmers) is the use of constants and enumerated types, both of which are supported by Pascal, C, and most other languages in some form. Constants allow the programmer to give a name to a number. In this way, the meaning of the number can be captured in the constant name, thus providing more information to the reader. Another advantage is that it is easy to change the value of the constant - you only need to change it at the definition.

Enumerated types allow the programmer to declare a variable with a limited range of values. These values can be numeric, or arbitrary strings. For example, if you want a variable to store a day of the week, instead of storing an integer value of 1 for Monday, you can store the more meaningful value `Monday`. Program status variables (i.e., state variables) and case statement indices should never be integer values. They should always have logical names, defined either through constants or enumerated types. You should carefully consider every number that you type in your code, and substitute a name for it unless you have a good reason not to do so.

The more that you limit the possible values for a variable (through the use of range and enumerated types), the less chance for error. Range and enumerated types also give error

checking compilers a chance to work in your favor, as well as supplying more information to the reader.

Records, Structures, Types and Typedefs

The most difficult part in writing many programs is defining the various record types (structures and unions in C). This is because the record definitions define the data structures for the program. You should give special care to naming both the record and its subfields, even more care than to naming single variables. Below is a Pascal example using records and enumerated types. Even without comments, the intention should be pretty clear. Note that I gave the record field the ``good" name at the expense of the name for the range and enumerated types, since the field name will be used a lot within the program code.

```
WeekdayType = (Mon, Tue, Wed, Thur, Fri, Sat, Sun);
MonthType = (Jan, Feb, Mar, Apr, May, Jun,
             Jul, Aug, Sep, Oct, Nov, Dec);
DayType = 1 .. 31;
Date = record
    month : MonthType;
    day : DayType;
    weekday : WeekdayType;
    year : integer
end;
```

Block Statements

Block statements (the group of statements within a begin-end pair in Pascal or between braces in C) deserve a comment. This comment should indicate what is happening in the block, such as the action for each pass through the loop, or what action is being executed in the else part of an if statement. In addition, at the end of the block, a comment should indicate what block is being closed. Here is an example in C.

```
while (listptr != NULL)
{ /* sum all values on the list */
    total += listptr->value;
    listptr = listptr->next;
} /* while (listptr != NULL) */
```