

Homework/Administrivia

- › **Read** §9.4.1–9.4.3 on two-dimensional arrays – *not covered in lecture*
- › **Exam 2 Bonus** Friday, November 21, in class – *details in previous slide deck*
- › **Read** the following from **Tarnoff-Ch13.pdf** (in Canvas under Files > Readings):
 - › §13.1 – Characteristics of the Memory Hierarchy
 - › §13.4 – Cache RAM
 - › §13.5 – Registers

Memory Characteristics (Review)

Memory technology	Typical access time	\$ per GiB in 2012
SRAM semiconductor memory	0.5–2.5 ns	\$500–\$1000
DRAM semiconductor memory	50–70 ns	\$10–\$20
Flash semiconductor memory	5,000–50,000 ns	\$0.75–\$1.00
Magnetic disk	5,000,000–20,000,000 ns	\$0.05–\$0.10

Table: Patterson & Hennessy, p. 378

Memory Hierarchy

- › A **memory hierarchy** is
 - › an organization of storage devices
 - › that takes advantage of the characteristics of different storage technologies
 - › to improve the overall performance of a computer system

Hard Drive vs. Main Memory

- › Accessing a hard drive is slow
- › Accessing main memory (DRAM) is significantly faster
- › **∴ Load programs and data from the hard drive into main memory as needed**
 - › Amount of RAM is usually much smaller than hard drive capacity – no problem
 - › Programs load the data they need, as they need it – rarely do they need *every* bit of data on the drive!

Main Memory vs. Cache Memory

- › Accessing a hard drive is slow
- › Accessing main memory (DRAM) is significantly faster
- › But accessing cache memory (SRAM) is significantly faster than accessing DRAM!
- › **∴ Use cache memory to make accesses to main memory appear faster (how?...)**
 - › Amount of cache memory is much smaller than the amount of main memory (*how to deal with this?...*)

Principle of Locality



- ▶ **Principle of Locality:**
 - ▶ Instructions that are executed within a short period of time tend to be close together in memory
 - ▶ Data that are accessed within a short period of time tend to be close together in memory
- ▶ Programs tend to exhibit both
 - ▶ **temporal locality:** if we access this instruction/data now, we are likely to need this same instruction/data again in the near future
 - ▶ **spatial locality:** if we access this instruction/data now, we are likely to need nearby instructions/data in the near future
- ▶ **Q.** Why is this true?

Cache: Concept



- ▶ Use the Principle of Locality to make main memory accesses appear faster:
- ▶ Place a small, fast SRAM (**cache**) between the processor and main memory




- ▶ Idea:
 - ▶ When the processor requests data, check if it's already in the cache
 - ▶ If not, copy it from main memory into the cache
 - ← Temporal locality suggests we'll need this data again soon
 - ▶ Also copy data in nearby memory locations
 - ← Spatial locality suggests we'll need surrounding data too

Cache: Hits and Misses



- ▶ Use the Principle of Locality to make main memory accesses appear faster:
- ▶ Place a small, fast SRAM (**cache**) between the processor and main memory




- ▶ Idea:
 - ▶ When the processor requests data, check if it's already in the cache
 - ▶ If the data is already in the cache, it is a **cache hit**
 - ▶ If the data is not already in the cache, it is a **cache miss**
 - ▶ Percentage of accesses that result in cache hits is the **hit rate** the **hit ratio**

Cache: Blocks & Cache Lines



- ▶ Divide main memory into **blocks** of a fixed size
 - ▶ Example below: 64 bytes (64 = 40h)
 - ▶ Note: each 64-byte group starts at a memory address that is divisible by 64 = 40h
- ▶ Cache will store copies of some of these blocks (the ones actively being used at a point in time)
 - ▶ When talking about blocks inside the cache, they are often called **cache lines**
- ▶ If *any* byte in a block is accessed, copy the **entire block** into the cache
 - ▶ E.g., access memory address 00000078h ⇒ copy block starting at address 00000040h



On an earlier slide, this is what was meant by "copy data in nearby memory locations"

A Concrete Example



```

Memory1
Address: 0x00401010
0x00401010 b0 00 fe c0 3c ff 72 fa *.p&cyru
0x00401018 6a 00 e8 0b 00 00 90 j.e....
0x00401020 90 90 90 90 cc cc cc ....llll

main.asm
.code
main PROC
    mov al, 0
    a: inc al
    cmp al, 255
    jb a
    push 0
    call ExitProcess
    nop
    nop
    nop
    nop
    main ENDP
  
```

- ▶ Suppose 64-byte cache lines
 - ▶ Note: main starts at address 00401010h, which is divisible by 64
- ▶ **Q.** Suppose data are copied into the cache when the processor fetches the first instruction. What other instructions will also be copied into the cache?
- ▶ **Q.** Suppose data are copied into the cache when the processor fetches the cmp instruction. What other instructions will also be copied into the cache?

Identifying Blocks (1)



- 0 = 0000 0000
 - 1 = 0000 0001
 - 2 = 0000 0010
 - 3 = 0000 0011
 - 4 = 0000 0100
 - 5 = 0000 0101
 - 6 = 0000 0110
 - 7 = 0000 0111
 - 8 = 0000 1000
 - 9 = 0000 1001
 - 10 = 0000 1010
 - 11 = 0000 1011
 - 12 = 0000 1100
- ▶ Suppose we have 4-byte blocks
 - ▶ Each block starts at a memory address divisible by 4
 - ▶ **The upper bits of all memory addresses within a block are the same!**
 - ▶ Only the lowest two bits are different (since 4 = 2²)
 - ▶ Idea:
 - ▶ Use the upper bits (that are the same within a block) to number that block
 - ▶ This is called the **block number**
 - ▶ Recall: You can divide by 2^k by right-shifting
 - ▶ Splitting the bits of a number into two groups gives the quotient and remainder when the number is divided by a power of 2
 - ▶ So the block number is the memory address divided by the block size

Identifying Blocks (2)



Suppose:

- 20-bit memory addresses (x86 real-address mode)
- 4-byte cache blocks

- $4 = 2^2$ bytes per block

- Q.** Given a memory address, what *bitwise operations* would you use to find

- the block number?
- the position of that byte within the block?

Memory address	Data	Block identification
0000 0000 0000 0000 00 00		Block 0
0000 0000 0000 0000 00 01		
0000 0000 0000 0000 00 10		
0000 0000 0000 0000 00 11		
0000 0000 0000 0000 01 00		Block 1
0000 0000 0000 0000 01 01		
0000 0000 0000 0000 01 10		
0000 0000 0000 0000 01 11		
0000 0000 0000 0000 10 00		Block 2
0000 0000 0000 0000 10 01		
0000 0000 0000 0000 10 10		
0000 0000 0000 0000 10 11		
...		
1111 1111 1111 1111 11 00		Block $2^{(20/2)} - 1 = 262,143$
1111 1111 1111 1111 11 01		
1111 1111 1111 1111 11 10		
1111 1111 1111 1111 11 11		

Each gray block represents an addressable memory location

Identifying Blocks (3)



- Q.** Suppose you have a 1 GB (2^{30} -byte) address space divided into 8-byte blocks.

- Note:* memory addresses are 30 bits

- (a) How many blocks are in the address space? (*Answer: 2 to what power?*)
- (b) How would a memory address be divided to identify a block and offset?

- A.** (a) $2^{30} \div 8 = 2^{30} \div 2^3 = 2^{27} = 134,217,728$

- (b) Memory address \rightarrow $a_{29} \ a_{28} \ a_{27} \ \dots \ a_4 \ a_3 \ a_2 \ a_1 \ a_0$

Bits identifying block
Bits identifying offset

Cache Entries



- A cache consists of several entries
- Number of entries is usually a power of 2

- Each entry contains (at least)

- the data from one block of memory
- a **tag** that identifies which block of memory is currently stored there

We defined a **cache line** to be synonymous with a block. The assigned reading defines a **line** to be an entire cache entry (tag+block), which is unusual.

TAG ₀	Block of words corresponding to TAG ₀	LINES
TAG ₁	Block of words corresponding to TAG ₁	
TAG ₂	Block of words corresponding to TAG ₂	
TAG ₃	Block of words corresponding to TAG ₃	
...	...	
TAG _{n-1}	Block of words corresponding to TAG _{n-1}	????

Cache Questions



- Remember: cache is significantly smaller than main memory

- Four questions in designing a cache system:

- block placement/mapping function** – where can a block of memory go in the cache?
- block identification** – how to determine if a block is in the cache, and where it is?
- replacement algorithm** – when the cache is full, what do you remove?
- write policy** – when the processor needs to *write* to memory, what happens?

Block Placement Policies



Direct-mapped

- For any block of memory, there's only one place where it can be stored
- Use (block number) mod (number of blocks) to determine which slot to store it in

Fully associative

- Any block of memory can be stored in any slot

Set-associative ← Most commonly used

- Cache entries are grouped together into *sets*
- The block number determines which *set* the block will be stored in
- But it can be stored in any slot within that set



Block Placement: Direct Mapping (1)



- In a **direct-mapped** cache, the block number determines where the block will be stored in the cache

- Number rows in the cache from 0; store in (block number) mod (number of rows)
- If there are 2^l blocks, the lowest l bits of the block number give the row number to store in
- tag = (block number) / (number of rows)
- Use the upper bits as a **tag** to identify which block is stored in this row


- Decompose a memory address into:



Block Placement: Direct Mapping (2)




- **Block Identification** – When a memory read occurs, determine if the corresponding block of memory is in the cache:



Block Placement: Fully Associative



- In a **fully associative** cache, any block can be stored in any row in the cache
- Use the entire block number for the tag




- To check if a block is in the cache, must check every row – it could be anywhere
- Wouldn't this be slow?
- No – Hardware can be designed to search all cache entries simultaneously (in parallel)
 - But this can be expensive

Block Placement: Set Associative



- In a **set associative** cache, rows are grouped together into *sets*
- If there are n entries in a set, the cache is said to be **n -way set associative**
- Block number determines which *set* the block will be stored in
- Decompose the memory address as before:
- Each *set* is a small associative cache (search like we did for a fully associative cache)



Cache Questions



- Remember: cache is significantly smaller than main memory
- Four questions in designing a cache system:
 - **block placement/mapping function** – where can a block of memory go in the cache?
 - **block identification** – how to determine if a block is in the cache, and where it is?
 - **replacement algorithm** – when the cache is full, what do you remove?
 - **write policy** – when the processor needs to *write* to memory, what happens?

Replacement Policy



- In a fully associative or set associative cache, a block could be placed in several possible places
- If the cache is full, one block must be removed (**evicted**) from the cache – which one?
- Choose a **replacement algorithm**:
 - **Least Recently Used (LRU)** ← Most commonly used
 - Replace the block that hasn't been read by the processor in the longest period of time
 - **First In First Out (FIFO)**
 - Replace the block that has been in the cache the longest
 - **Least Frequently Used (LFU)**
 - Replace the block with the fewest hits since being loaded into the cache
 - **Random**
 - Randomly select a block to replace

Cache Questions



- Remember: cache is significantly smaller than main memory
- Four questions in designing a cache system:
 - **block placement/mapping function** – where can a block of memory go in the cache?
 - **block identification** – how to determine if a block is in the cache, and where it is?
 - **replacement algorithm** – when the cache is full, what do you remove?
 - **write policy** – when the processor needs to *write* to memory, what happens?

Cache Write Policy



- Problem: There are now *two* copies of data – one in the cache, one in main memory
- What happens when the processor wants to *write* to memory?




- Choose a **cache write policy**:
 - Write-through**: when the processor writes, the cache updates itself *and* main memory
 - Write-back**: only update the cache; update main memory when the block is removed from the cache


More Advanced Cache Designs



- If instructions (code) and data are both stored in the same cache, it is a **unified** cache
- Some processors use **split caches** – separate caches for code (instructions) and data
 - Instruction cache and data cache often called **I-cache** and **D-cache**, respectively



- Modern processors use two or three levels of cache memory (**multilevel caches**)



Example: Intel “Sandy Bridge”



- Intel’s Sandy Bridge microarchitecture – released 2011, used in 2nd generation Core i3, i5, i7
- L1 Cache: 64 KB per core, split, 8-way set associative, write-back, 64-byte cache lines
- L2 Cache: 256 KB per core, unified, 8-way set associative, write-back, 64-byte cache lines
- L3 Cache: size varies (1–20 MB), 64-byte cache lines

From the Intel® 64 and IA-32 Architectures Optimization Reference Manual:

Table 2-10. Cache Parameters

Level	Capacity	Associativity (ways)	Line Size (bytes)	Write Update Policy
L1 Data	32 KB	8	64	Writeback
Instruction	32 KB	8	N/A	N/A
L2 (Unified)	256 KB	8	64	Writeback
Third Level (LLC)	Varies, query CPUID leaf 4	Varies with cache size	64	Writeback

Source: <http://www.intel.com/content/dam/www/public/us/en/documents/manuals/64-ia-32-architectures-optimization-manual.pdf>

Does cache memory matter?

